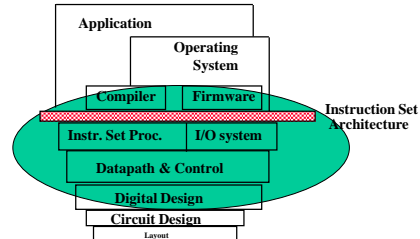
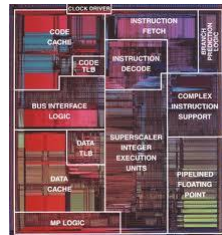


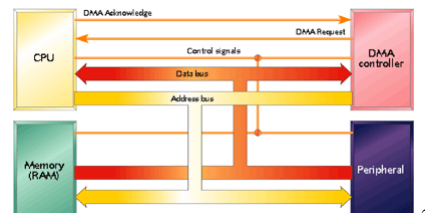
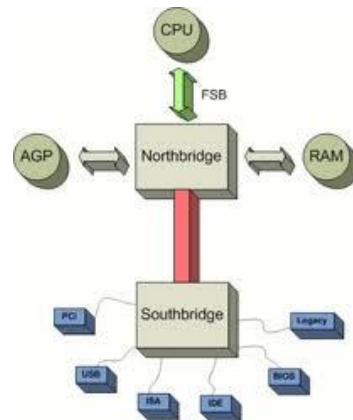
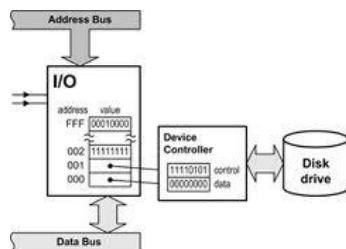
CS/SE 3340

Computer Architecture



Input & Output

Adapted from slides by Profs. D. Patterson and J. Hennessey



Why Input & Output?

- So far we have learnt how the processor process data in memory
 - Huge array of bytes
- But *how* does the processor process information from **outside world**?
 - Get data from a keyboard or a sensor
 - Display results of processing on screen (monitor) or paper (printer)?
 - Communicate with other system via networks?
- The processor needs to work with input & output (I/O) devices

3

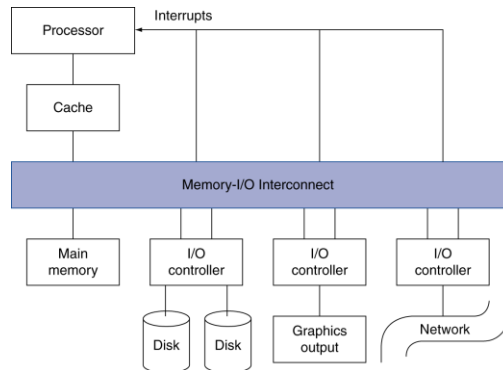
From I/O Channels to Device Drivers

- In the beginning **I/O channels** provides means for the CPU to work with physical I/O devices
 - I/O hardware understands commands from CPU
 - The CPU has special instructions to pass commands to I/O hardware (e.g. *Channel Control Words, CCW*, in IBM mainframes)
- Modern Operating Systems unifies I/O operations under **device drivers**
 - Layer of **abstraction** that hides I/O device specific details

4

I/O Devices

- I/O devices can be characterized by
 - Function: *input, output, storage*
 - Interface with the CPU: *serial, parallel*
 - Unit of data transfer: *character, block*
- I/O interconnect allows for data exchange among components



5

Interconnecting Components

- Need interconnections between
 - CPU, memory, I/O controllers
- **Bus**: shared communication channel
 - Parallel set of wires for data and synchronization of data transfer
 - Can become a bottleneck
- Performance limited by physical factors
 - Wire length, number of connections
- More recent alternative: high-speed serial connections with switches
 - Like networks

6

Bus Types

- Processor-Memory buses
 - Short, high speed
 - Design is matched to memory organization
- I/O buses
 - Longer, allowing multiple connections
 - Specified by standards for interoperability
 - Connect to processor-memory bus through a bridge
 - e.g. Universal Serial Bus (USB)

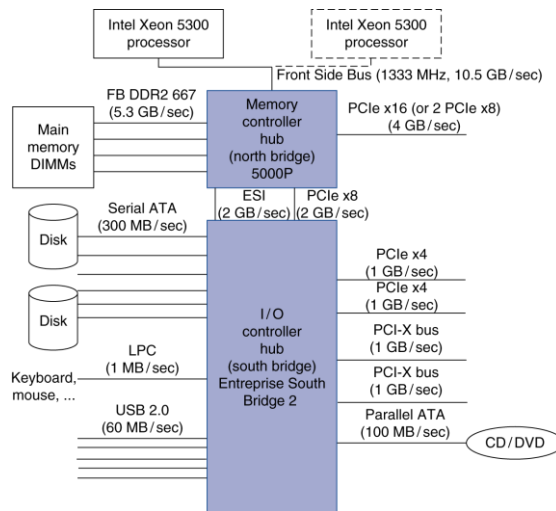


7

I/O Bus Examples

	Firewire	USB 2.0, USB 3.0	PCI Express	Serial ATA	Serial Attached SCSI
Intended use	External	External	Internal	Internal	External
Devices per channel	63	127	1	1	4
Data width	4	2	2/lane	4	4
Peak bandwidth	50MB/s or 100MB/s	0.2MB/s, 1.5MB/s, or 60MB/s 600MB/s	250MB/s/lane 1×, 2×, 4×, 8×, 16×, 32×	300MB/s	300MB/s
Hot pluggable	Yes	Yes	Depends	Yes	Yes
Max length	4.5m	5m	0.5m	1m	8m
Standard	IEEE 1394	USB Implementers Forum	PCI-SIG	SATA-IO	INCITS TC T10 8

Typical x86 PC I/O System



9

I/O Management

- In modern computer systems I/O operations are mediated by the OS
- I/O operations
 - Multiple programs share I/O resources
 - Need protection and scheduling
 - I/O causes asynchronous interrupts
 - Same mechanism as exceptions
 - I/O programming is fiddly
 - OS provides abstractions to programs

10

I/O Registers

- I/O devices are managed by I/O controller hardware
 - Transfers data to/from device
 - Abstracted as three sets of registers
- **Command registers**
 - Cause device to do something
- **Status registers**
 - Indicate what the device is doing and occurrence of errors
- **Data registers**
 - Write: transfer data to a device
 - Read: transfer data from a device

11

Where To Locate I/O Registers?

- *Memory mapped I/O*
 - Registers are addressed in same space as memory
 - Address decoder distinguishes between them
 - OS uses address translation mechanism to make them only accessible to kernel
 - Example: MIPS
- *I/O instructions and I/O ports*
 - Separate instructions to access I/O ports
 - Can only be executed in kernel mode
 - Example: x86 (IN and OUT instructions on I/O segment)

12

When to Perform I/O Operations ? Easy Way - Polling

- Periodically check I/O status register
 - If device ready, do operation
 - If error, take action
- Common in small or low-performance real-time embedded systems
 - Predictable timing
 - Low hardware cost
- In other systems, wastes CPU time

13

When to Perform I/O Operations? Better Way - Interrupts

- When a device is ready or error occurs
 - Controller *interrupts* CPU
- *Interrupt* is like an exception
 - But not synchronized to instruction execution
 - Can invoke handler between instructions
 - Cause information often identifies the interrupting device
- Priority interrupts
 - Devices needing more urgent attention get higher priority
 - A higher priority I/O device can interrupt handler for a lower priority I/O device

14

I/O Data Transfer – How?

- *Programmed I/O*
 - CPU transfers data between memory and I/O data registers
 - Time consuming for high-speed devices
- *Direct Memory Access (DMA)*
 - OS provides starting address in memory
 - I/O controller transfers to/from memory autonomously
 - Controller interrupts on completion or error

15

MIPS I/O

- MIPS supports *interrupts* and *DMA*
- MIPS CPU communicates with I/O devices using *memory-mapped* input/output (I/O)
- With memory mapped I/O, the I/O devices appear to be primary memory locations
 - There is no need for specific (additional) instructions to communicate with I/O devices
 - No need for I/O ports either
- What existing instructions can be used?
 - *Load & Store*

16

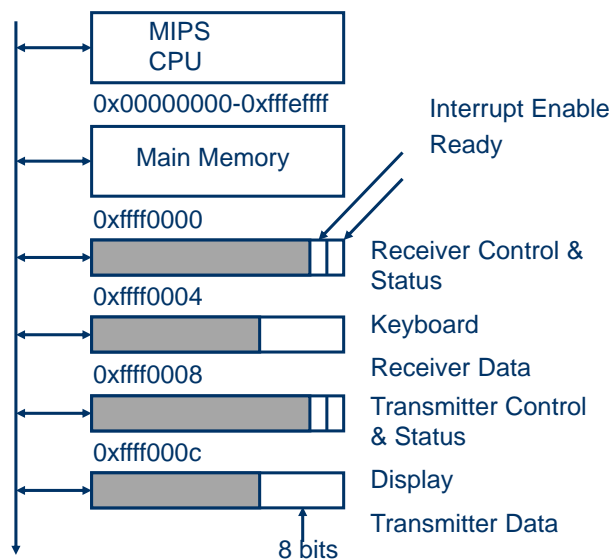
MIPS- Memory-Mapped I/O

- The MIPS I/O address space **0xffff0000** to **0xffffffff** is reserved for memory-mapped I/O
 - How large is the I/O address space?
- Any address in this space can be mapped to I/O registers (command, status, data) of an I/O device
- I/O is performed using the regular *load* and *store* instructions on these addresses
 - e.g.

```
li    $t0, 0xffff0004
lw    $s1, ($t0)
```

17

MIPS I/O Example



18

Keyboard Controller – C Language

Keyboard Control	0xFFFF0000	2 bits
Keyboard Data	0xFFFF0004	8 bits
Display Control	0xFFFF0008	2 bits
Display Data	0xFFFF000C	8 bits

```
char read_ch (void) {
    volatile unsigned int *recv;
    recv = 0xFFFF0000;
    while ((recv[0] & 1) == 0) /* - */;
    return recv[1] & 0x00FF;
}
```

19

Keyboard Controller – MIPS Assembly

```
.text
.globl read_ch
read_ch:
    subu    $sp, $sp, 4 # "push" stack
    sw      $t0, 0($sp) # save temporaries
    li      $t0, 0xffff0000 # RCR address
recv_rda:
    lw      $v0, 0($t0) # read status from RCR
    andi    $v0, $v0, 0x0001 # check if Ready
    beq     $v0, $zero, recv_rda # wait until
Ready is "one"
    lw      $v0, 4($t0) # read data from RDR
    lw      $t0, 0($sp) # temporaries
    addiu   $sp, $sp, 4 # "pop" stack
    jr      $ra
```

20

Display Controller – C Language

Keyboard Control	0xFFFF0000	2 bits
Keyboard Data	0xFFFF0004	8 bits
Display Control	0xFFFF0008	2 bits
Display Data	0xFFFF000C	8 bits

```
void write_ch (char c){
    volatile unsigned int  *txmit;

    txmit = 0xFFFF0008; /* base */
    while ((txmit[0] & 1) == 0) /* - */;
    txmit[1] = c;
}
```

21

Display Controller – MIPS Assembly

```
.text
.globl write_ch      # ARG0 = $a0 = character
write_ch:
    subu    $sp, $sp, 8 # "push" stack
    sw      $t1, 4($sp) # save temporaries
    sw      $t0, 0($sp) # save temporaries
    li      $t0, 0xffff0008 # TCR address
write_tbe:
    lw      $t1, 0($t0) # transmitter status port
    andi    $t1, $t1, 0x0001 # TDR is empty?
    beq     $t1, $zero, write_tbe # is READY "zero"
    sw      $a0, 4($t0) # write data-byte to TDR
    lw      $t1, 4($sp) # temporaries
    lw      $t0, 0($sp) # temporaries
    addiu   $sp, $sp, 8 # "pop" stack
    jr      $ra
```

22

Summary

- The CPU processes data from outside world via input and output operations
- Device drivers are abstraction that hides I/O device specifics
- *Where*: memory mapped or I/O ports
- *When*: CPU can either poll for data, or be interrupted when data are ready
- *How*: programmed I/O or DMA based
- Memory-mapped I/O eliminates the need for I/O specific instructions and I/O ports

23