
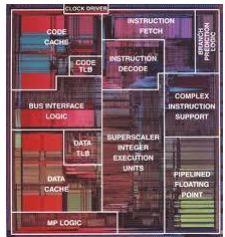




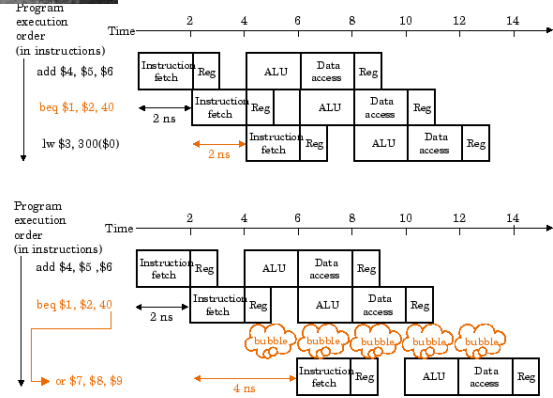
# CS/SE 3340

## Computer Architecture



### Pipelining and Hazards

Adapted from slides by Profs. D. Patterson and J. Hennessey



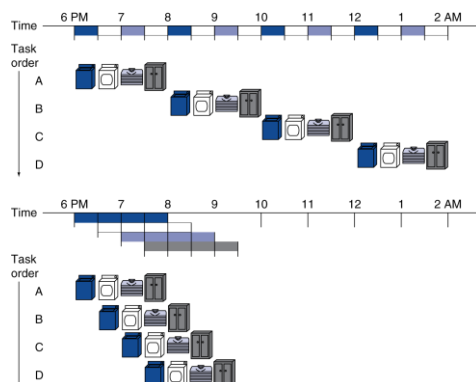
## Performance Issues

- Longest delay determines clock period
  - **Critical path**: load instruction
  - *Instruction memory* → *register file* → *ALU* → *data memory* → *register file*
- Not feasible to vary period for different instructions
- Violates design principle
  - Making the common case fast
- We will improve performance by **pipelining**

3

## Pipelining Analogy

- Pipelined laundry: *overlapping execution*
  - (temporal) Parallelism improves performance



■ Four loads:

■ Speedup  
 $= 8/3.5 = 2.3$

■ Non-stop:

■ Speedup  
 $= 2n/0.5n + 1.5 \approx 4$   
 $= \text{number of stages}$

4

## MIPS Pipeline

- Five stages, one step per stage
  1. **IF**: Instruction fetch from memory
  2. **ID**: Instruction decode & register read
  3. **EX**: Execute operation or calculate address
  4. **MEM**: Access memory operand
  5. **WB**: Write result back to register

5

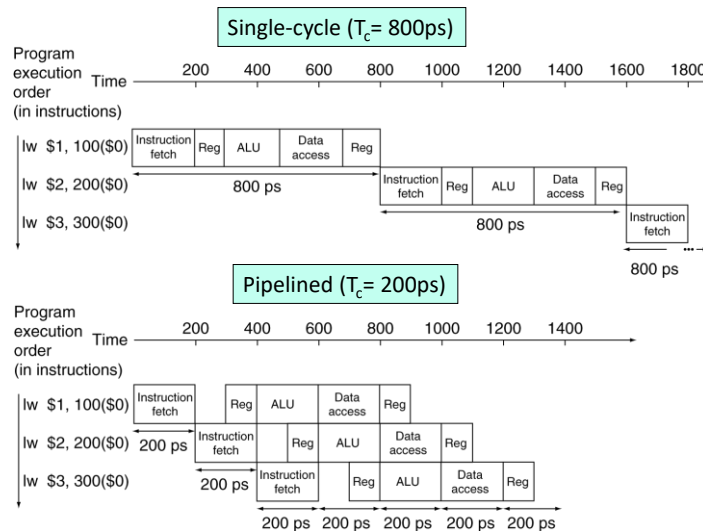
## Pipeline Performance

- Assume time for stages is
  - 100ps for register read or write
  - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

| Instr    | Instr fetch | Register read | ALU op | Memory access | Register write | Total time |
|----------|-------------|---------------|--------|---------------|----------------|------------|
| lw       | 200ps       | 100 ps        | 200ps  | 200ps         | 100 ps         | 800ps      |
| sw       | 200ps       | 100 ps        | 200ps  | 200ps         |                | 700ps      |
| R-format | 200ps       | 100 ps        | 200ps  |               | 100 ps         | 600ps      |
| beq      | 200ps       | 100 ps        | 200ps  |               |                | 500ps      |

6

# Pipeline Performance



7

# Pipeline Speedup

- If all stages are balanced
  - i.e., all take the same time
  - Time between instructions<sub>pipelined</sub> =  $\frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of stages}}$
- If not balanced, speedup is less
- Speedup due to increased throughput
  - Latency (time for each instruction) *does not* decrease

8

## Pipelining and ISA Design

- MIPS ISA designed for *pipelining*
  - All instructions are 32-bits
    - Easier to fetch and decode in one cycle
    - c.f. x86: 1- to 17-byte instructions
  - Few and regular instruction formats
    - Can decode and read registers in one step
  - Load/store addressing
    - Can calculate address in 3<sup>rd</sup> stage, access memory in 4<sup>th</sup> stage
  - Alignment of memory operands
    - Memory access takes only one cycle

9

## Hazards

- Situations that prevent starting the next instruction in the next cycle
- *Structure hazards*
  - A required resource is busy
- *Data hazard*
  - Need to wait for previous instruction to complete its data read/write
- *Control hazard*
  - Deciding on control action depends on previous instruction

10

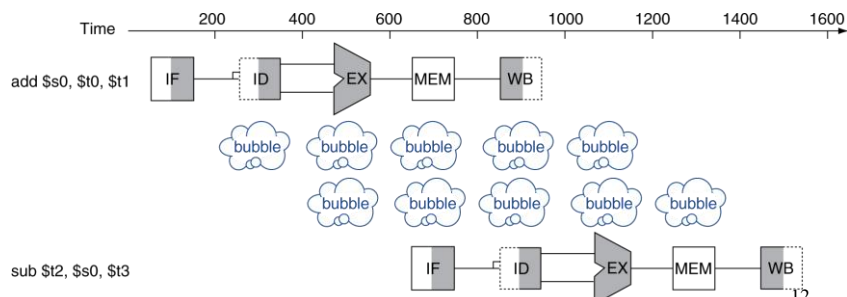
## Structure Hazards

- Conflict for use of a resource
- In MIPS pipeline with a single memory
  - Load/store requires data access
  - Instruction fetch would have to *stall* for that cycle
    - Would cause a pipeline “bubble”
- Hence, pipelined datapaths require separate instruction/data memories
  - Or separate instruction/data caches

11

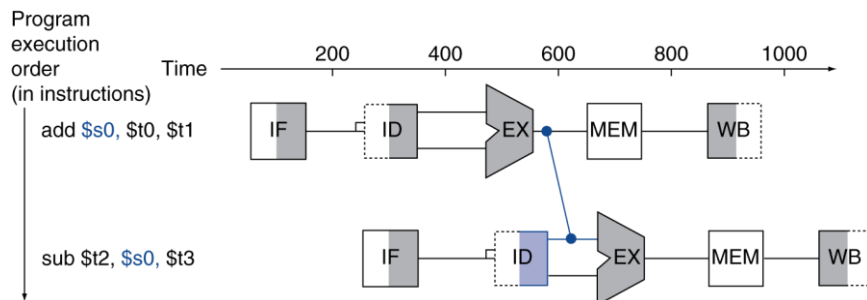
## Data Hazards

- An instruction depends on completion of data access by a previous instruction
  - add **\$s0**, \$t0, \$t1
  - sub \$t2, **\$s0**, \$t3



## Forwarding (aka Bypassing)

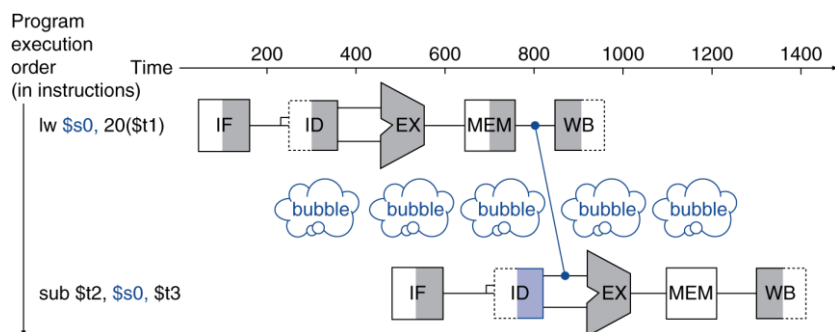
- Use result when it is computed
  - Don't wait for it to be stored in a register
  - Requires extra connections in the datapath



13

## Load-Use Data Hazard

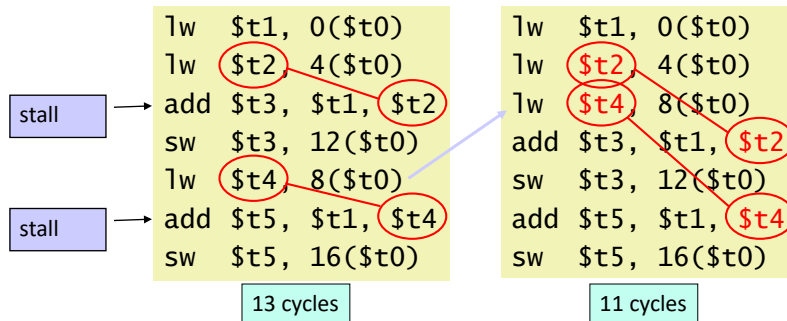
- Can't always avoid stalls by forwarding
  - If value not computed when needed
  - Can't forward backward in time!



14

## Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for  $A = B + E$ ;  $C = B + F$ ;



15

## Control Hazards

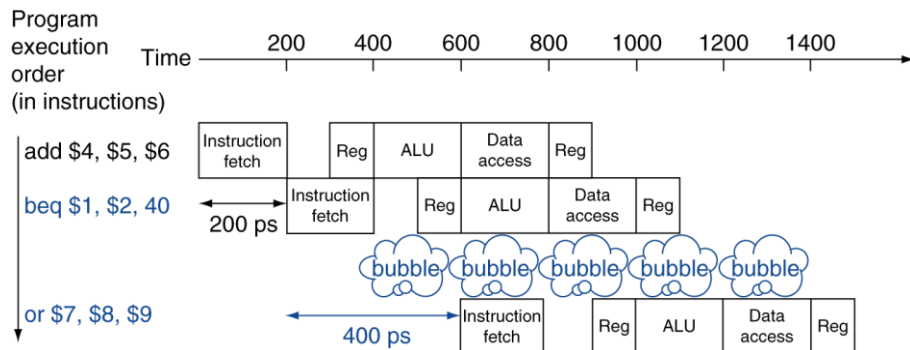
- Branch determines flow of control
  - Fetching next instruction depends on branch outcome
  - Pipeline can't always fetch correct instruction
    - Still working on ID stage of branch
- In MIPS pipeline
  - Need to compare registers and compute target early in the pipeline
  - Add hardware to do it in ID stage

16



## Stall on Branch

- Wait until branch outcome determined before fetching next instruction



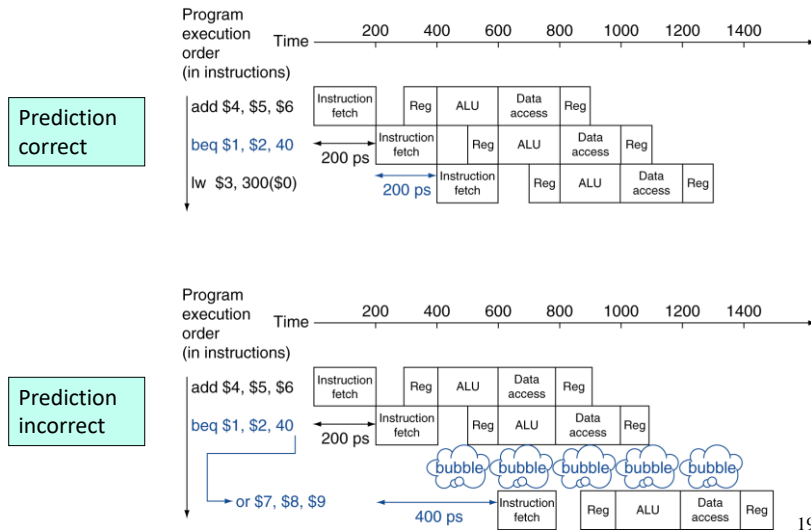
17

## Branch Prediction

- Longer pipelines can't readily determine branch outcome early
  - Stall penalty becomes unacceptable
- Predict outcome of branch
  - Only stall if prediction is wrong
- In MIPS pipeline
  - Can predict branches not taken
  - Fetch instruction after branch, with no delay

18

## MIPS with Predict: Not Taken



19

## More-Realistic Branch Prediction

- *Static branch prediction*
  - Based on typical branch behavior
  - Example: loop and if-statement branches
    - Predict backward branches taken
    - Predict forward branches not taken
- *Dynamic branch prediction*
  - Hardware measures actual branch behavior
    - e.g., record recent history of each branch
  - Assume future behavior will continue the trend
    - When wrong, stall while re-fetching, and update history

20

## Summary

- Modern processors take advantages of the “pipelined effect” to achieve higher performance
  - Instruction execution is divide into stages (ala assembly line)
  - If all stages are balanced at all time, an n-times speed up can be achieved (n: number of stages)
- Hazards cause stages of instruction execution pipeline to be stalled
- Various techniques have been developed to deal with hazards

21