





CS/SE 3340

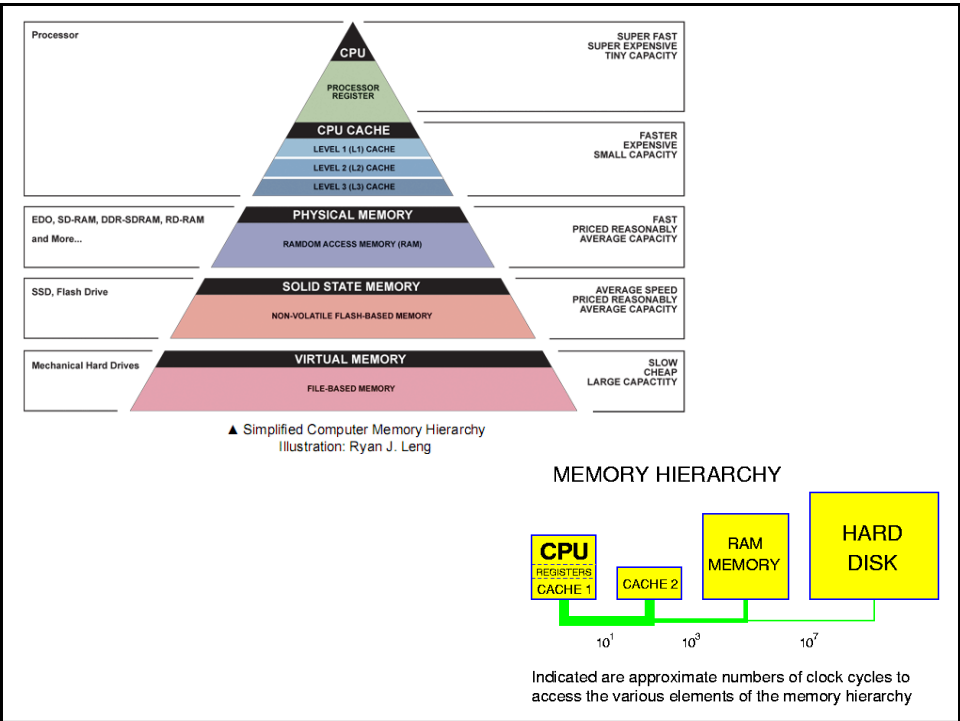
Computer Architecture



Memory Hierarchy – An Overview


Adapted from slides by Profs. M. Irwin, D. Patterson and J. Hennessey

1



2

Major Components of a Computer




3

3

The “Memory Wall”

- Processor v.s. main memory speed disparity continues to grow



Good **memory hierarchy** design is increasingly important to overall performance of a computer system!

4

4

The Memory Hierarchy Goal

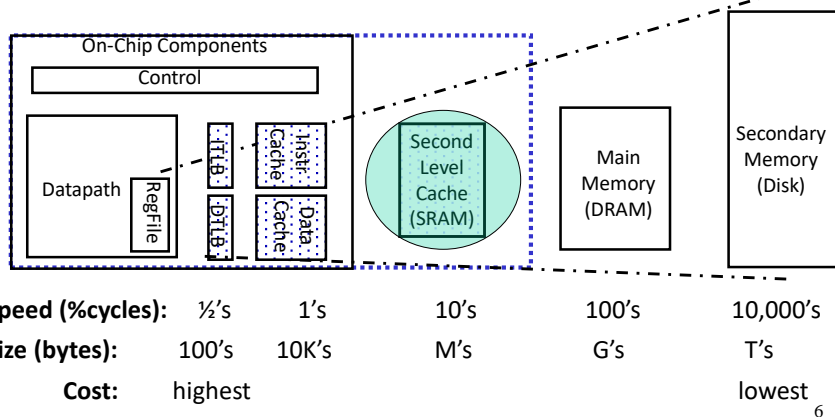
- Due to cost constraints
 - **Fast** memories are *small* (static RAM)
 - **Large** memories are *slow* (dynamic RAM)
- How do we create a memory system that gives the illusion of being *large*, *cheap* and *fast* (most of the time)?
 - With *hierarchy*
 - With *locality*

5

5

A Typical Memory Hierarchy

- Take advantage of the principle of locality to present the user with as much memory as is available in the *cheapest* technology at the speed offered by the *fastest* technology



6

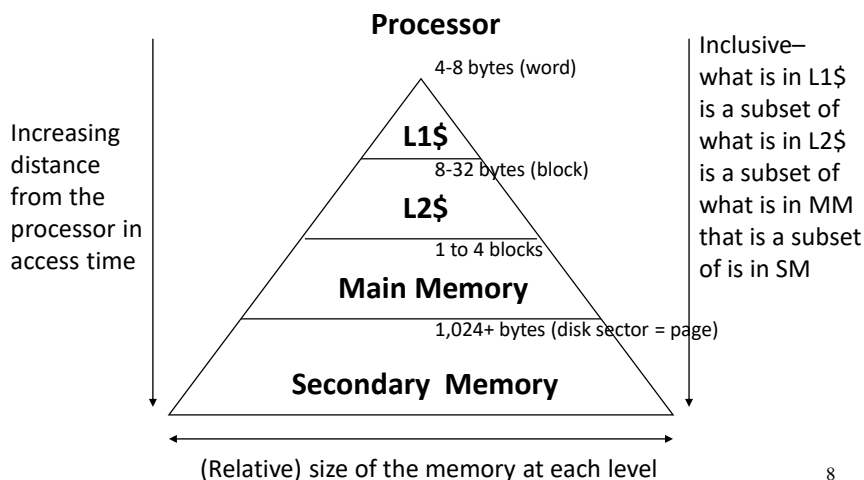
Principle of Locality

- Programs access a small proportion of their address space at any time
- *Temporal locality*
 - Items accessed recently are likely to be accessed again soon, e.g. instructions in a loop, induction variables
 - ⇒ Keep **most recently accessed** data items closer to the processor
- *Spatial locality*
 - Items near those accessed recently are likely to be accessed soon, e.g. array data
 - ⇒ Move blocks consisting of **contiguous words** closer to the processor

7

7

Characteristics of the Memory Hierarchy



8

8

How is the Hierarchy Managed?

- registers ↔ memory
 - by compiler (programmer?)
- cache ↔ main memory
 - by the cache controller hardware
- main memory ↔ disks
 - by the operating system (virtual memory)
 - virtual to physical address mapping assisted by the hardware (Translation Lookaside Buffer, TLB) -> next session's topic
 - by the programmer (files)

9

9


Cache Terminology

- **Capacity (C):**
 - number of data bytes in cache
- **Block (aka line) size (b):**
 - bytes of data brought into cache at once, unit of data transfer
- **Number of blocks ($B = C/b$):**
 - number of cache blocks in cache: $B = C/b$
- **Degree of associativity (N):**
 - number of cache blocks in a set
- **Number of sets ($S = B/N$):**
 - each memory address maps to exactly one cache set

10

10

Memory Hierarchy Levels




- **Block (aka line):** unit of copying
 - May be multiple words
- If accessed data is present in upper level
 - **Hit:** access satisfied by upper level
 - Hit ratio: hits/accesses
- If accessed data is absent
 - **Miss:** block copied from lower level
 - Time taken: miss penalty
 - Miss ratio: misses/accesses
= 1 – hit ratio
 - Then accessed data supplied from upper level

11

11

Cache Memory

- Cache memory
 - The level of the memory hierarchy closest to the CPU
- Given accesses X_1, \dots, X_{n-1}, X_n



a. Before the reference to X_n

b. After the reference to X_n

- How do we know if the data is present?
- Where do we look?

12

12

Where To Find Data on Cache?


- Cache organized into S sets
 - Each memory address maps to exactly one set
- Cache is categorized by # of blocks in a set:
 - **Direct mapped**
 - 1 block per set
 - **N-way set associative**
 - N blocks per set
 - **Fully associative**
 - all cache blocks in 1 set

13

13

Where – Direct Mapped Cache

- Location determined by address
- Direct mapped: only one choice
 - (Block address) modulo (#Blocks in cache)



- #Blocks is a power of 2 – *Why?*
- Use low-order address bits

14

14

How – Tags and Valid Bits

- How do we know which particular block is stored in a cache location?
 - Store block address as well as the data
 - Actually, only need the high-order bits
 - Called the **tag**
- What if there is no data in a location?
 - **Valid bit (V)**: 1 = present, 0 = not present
 - Initially 0

15

15

Cache Example

- 8-blocks, 1 word/block, direct mapped
- Initial state

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

16

16

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Miss	110

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

17

17

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
26	11 010	Miss	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

18

18

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Hit	110
26	11 010	Hit	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

19

19

Cache Example

Word addr	Binary addr	Hit/miss	Cache block
16	10 000	Miss	000
3	00 011	Miss	011
16	10 000	Hit	000

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

20

20

Cache Example

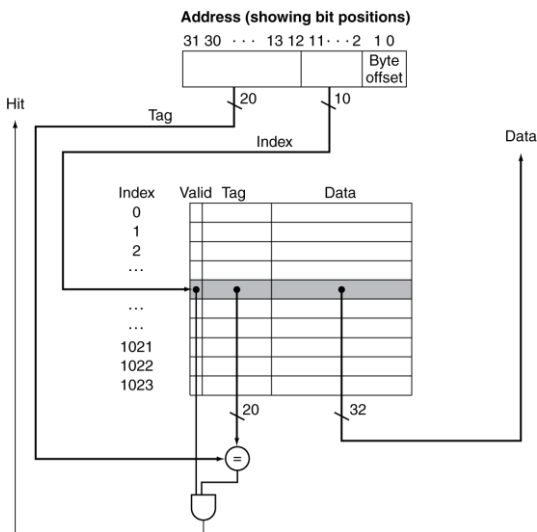
Word addr	Binary addr	Hit/miss	Cache block
18	10 010	Miss	010

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

21

21


Address Subdivision



22

22

Example: Larger Block Size



- 64 blocks, 16 bytes/block
 - To what block number does address 1200 map?
- Block address = $\lfloor 1200/16 \rfloor = 75$
- Block number = 75 modulo 64 = 11

23

23

Block Size Considerations

- Larger blocks should *reduce miss rate*
 - Due to spatial locality
- But in a fixed-sized cache
 - Larger blocks \Rightarrow fewer of them
 - More competition \Rightarrow *increased miss rate*
 - Larger blocks \Rightarrow pollution
- Larger miss penalty
 - Can override benefit of reduced miss rate
 - Early restart and critical-word-first can help

24

24

Cache Misses

- On cache hit, CPU proceeds normally
- On cache miss
 - Stall the CPU pipeline
 - Fetch block from next level of hierarchy
 - Instruction cache miss
 - Restart instruction fetch
 - Data cache miss
 - Complete data access

25

25

Write-Through

- On data-write hit, could just update the block in cache
 - But then cache and memory would be inconsistent
- *Write through*: also update memory
- But makes writes take longer
 - e.g., if base CPI = 1, 10% of instructions are stores, write to memory takes 100 cycles
 - Effective CPI = $1 + 0.1 \times 100 = 11$
- Solution: *write buffer*
 - Holds data waiting to be written to memory
 - CPU continues immediately
 - Only stalls on write if write buffer is already full

26

26

Write-Back

- On data-write hit, just update the block in cache
 - Keep track of whether each block is *dirty*
- When a *dirty* block is replaced
 - Write it *back* to memory
 - Can use a write buffer to allow replacing block to be read first

27

27

Summary

- Principle of **locality** (data access) in computer systems
 - *Temporal*
 - *Spatial*
- Memory hierarchy exploits this principle to present the processor with as much memory as is available in the cheapest technology at the speed offered by the fastest technology
- **Cache**: bring often used data closer to the CPU
- **Write access** requires special consideration to maintain data consistency: *write-through & write-back*

28

28